

PYDI SIDDARDHA

Snowflake Developer | Data Engineer | Cloud Data Platforms

+91 9945034950 • pydi.siddardha@gmail.com • linkedin.com/in/pydi11192307 • Bengaluru, IN

PROFESSIONAL SUMMARY

Results-driven Snowflake Developer and Data Engineer with 5+ years of experience architecting cloud-native data platforms at Fortune 500 clients (CVS Health, Morgan Stanley). Expert in end-to-end ELT pipeline design using Snowflake, DBT, Airflow, and PySpark. Proven track record migrating legacy Informatica/SSIS workflows to Snowflake with measurable performance gains, reducing pipeline runtimes and cutting infrastructure costs. Deep domain expertise in Healthcare (HIPAA) and Financial Services (MiFID II, SEC compliance). Immediate value-add for teams seeking to modernize data infrastructure on AWS/Azure.

CORE COMPETENCIES

Snowflake (Expert) • DBT (Advanced) • Apache Airflow • AWS S3 • Azure Databricks • PySpark • Matillion • Python • SQL • Medallion Architecture • SCD Type 1 & 2 • Snowpipe • CI/CD (GitHub) • ETL/ELT Migration • Data Modeling • HIPAA • Informatica PowerCenter • Teradata • Agile / Jira

TECHNICAL SKILLS

Cloud & Warehousing: Snowflake, AWS (S3, Glue), Azure (Databricks, ADF), Google BigQuery

Data Engineering: DBT, Apache Airflow, Matillion, Snowpipe, PySpark, Informatica PowerCenter

Databases: SQL Server, Teradata, MySQL

Languages: SQL (Advanced), Python, PySpark, Shell Scripting

Patterns & Architectures: Medallion (Bronze/Silver/Gold), Star/Snowflake Schema, SCD Type 1 & 2, Kimball

DevOps & Tooling: GitHub, Azure Repos, CI/CD, Jira, Confluence, Agile/Scrum

PROFESSIONAL EXPERIENCE

Capgemini — Snowflake Developer / Data Engineer

Mar 2021 – Present | Bengaluru, IN

Project: Pharmacy Data Analytics Platform (CVS Health – Caremark) | Sep 2022 – Present

Stack: Snowflake, DBT, Matillion, AWS S3, Airflow, Python, SQL Server, GitHub

Domain: Healthcare / Pharma — HIPAA-compliant claims & prescription analytics

Key Achievements:

- Led end-to-end migration of 40+ Informatica PowerCenter and SSIS pipelines to Matillion on Snowflake, reducing average pipeline runtime by ~35% and eliminating legacy licensing costs.
- Architected Medallion (Bronze → Silver → Gold) data lake on Snowflake for pharmacy claims, prescription, and drug pricing data from Oracle, SQL Server, and AWS S3 — enabling HIPAA-compliant reporting for 10M+ patient records.
- Implemented SCD Type 1 & Type 2 logic using DBT snapshots for historical patient drug adherence tracking, cutting manual reconciliation effort by 50% through automated DBT data quality tests.
- Configured Snowpipe for near-real-time continuous ingestion of JSON and Parquet files from S3, reducing data latency from batch (nightly) to under 15 minutes for fraud detection and cost-optimization alerts.
- Orchestrated multi-step Airflow DAGs with SLA enforcement, retry logic, and alerting for 20+ daily pipeline runs — achieving 99.6% on-time SLA adherence across production workloads.
- Optimized critical Snowflake queries by refactoring correlated sub-queries, applying cluster keys, and leveraging result caching — delivering 3–5x query speedup on high-frequency BI dashboards.
- Built and maintained 30+ incremental DBT models with custom macros, hooks, and tests, providing finance and clinical analytics teams with certified, reusable datasets.
- Executed source-to-target reconciliation across SQL Server/Synapse and Snowflake; achieved 100% row-count and hash-match parity across all migrated datasets.

Project: Trade Surveillance Data Platform (Morgan Stanley) | Mar 2021 – Sep 2022

Stack: Snowflake, Azure Databricks, PySpark, Teradata, AWS S3, GitHub, Jira

Domain: Investment Banking / FinTech — MiFID II & SEC compliance, market abuse detection

Key Achievements:

- Played a key role in migrating legacy Teradata data objects (tables, views, stored procedures) to Snowflake, ensuring functional parity and delivering 2x improvement in ad-hoc query performance for trading desks.
- Designed Star Schema dimensional models in Snowflake to support high-speed ad-hoc reporting for traders and risk managers — reducing report generation time from 8 min to under 90 seconds.
- Developed complex SQL stored procedures and views to detect insider trading and spoofing patterns, directly supporting MiFID II and SEC regulatory compliance reporting.

- Reduced storage costs by 28% by auditing Azure Databricks workloads with PySpark — identifying 150+ unused high-cost tables, removing duplicates, and applying Z-Ordering for efficient data layout.
- Built and deployed 10+ PySpark optimization scripts in Databricks notebooks; managed version control via GitHub with structured branching aligned to client release environments.
- Designed and maintained Snowflake objects (tables, views, external stages, file formats, sequences) for scalable ingestion of real-time trade event streams.
- Tracked and delivered 30+ Jira optimization tasks for table cleanup and performance improvements in collaboration with Morgan Stanley compliance and engineering stakeholders.

KEY ACHIEVEMENTS & IMPACT

- Automated 50% of manual data monitoring tasks using Python scripts and Snowflake Tasks, freeing ~10 hrs/week of analyst time.
- Managed and processed multi-terabyte datasets (structured + semi-structured: JSON, Parquet, Avro) across Snowflake and Databricks.
- Consistently delivered in Agile sprint cycles — participated in sprint planning, code reviews, and CI/CD release cycles across 5+ releases per quarter.
- Domain expertise across two regulated industries: Healthcare (HIPAA, PHI) and Financial Services (MiFID II, SEC) — a rare combination for Snowflake developers.

EDUCATION

Bachelor of Technology (B.Tech) — Computer Science Engineering

Acharya Nagarjuna University | 2014

DECLARATION

I hereby declare that the above furnished details are true and accurate to the best of my knowledge.

(Pydi Siddardha)